


UNIT 2: UNIX+OS CONCEPTS

OVERALL STRUCTURE — SYSTEM VS FUNCTION CALLS




	System Call	Function Call
Description	Application calls to functions provided by OS kernel	Program calls to predefined functions (eg: defined in library)
Behavior	Causes switches from user to kernel (space/mode)	Cause no space/mode switches
Speed	Slow	Fast


PROCESS MANAGEMENT (Time-sharing: Processes can "share" CPU time)

Context-switch: The process of storing the execution context of a process; and restoring the execution context of a new process

Context-switching causes overload.



As a flowchart, we have:



FILE/DIRECTORY MANAGEMENT

In Unix, everything is a file.

Root directory: Top-most dir. "/"


Home directory: Directory when you log in

begin at root

Absolute path: /x/y/z Relative path: x/y/z

THE MOUNT COMMAND

Here, we mount in "/mnt",



SYNTAX

`s=mkdir(name, mode)` create dir
`s=rmdir(name)` remove dir
`s=link(name1, name2)` Create new name2=name1
`s=unlink(name)` remove dir entry
`s=mount(special, name, flag)` mount
`s=umount(special)` unmount

Remark:


Pipe ("|") sends the output of one process as the input for another.

UNIT 3: FILE I/O (PART 1)

WEEK 2 Shun/335 (@shun4mick)

FILE I/O (BUFFERED VS UNBUFFERED)

Disk structure:



Buffered (standard) I/Os: Functions accumulate results in intermediate buffers, not making system calls each time (e.g. `fread/fwrite`)

Unbuffered I/Os: Functions invoke system calls to the kernel each time (e.g. `read()`, `write()` in Unix)

Of course, buffered is faster due to no system calls. Unbuffered would require updating process, OS kernel, and storage as needed. Even with buffer cache in OS kernel, it takes a long time. Besides, buffer cache needs to be maintained carefully alongside storage when `write()` is called.

SYNTAX — FILE DESCRIPTORS


File descriptor: A nonnegative integer $\in [0, \text{OPEN_MAX}-1]$, where $\text{OPEN_MAX} = \text{max files a process can open at once}$

↳ They are per-process, so diff processes may share file descriptors

In `<unistd.h>`, we have `0=STDIN_FILENO`, `1=STDOUT_FILENO`, `2=STDERR_FILENO` by POSIX.1 standard. More are stored in a file table.

```
char buf[100];
E.g. while (n=read(STDIN_FILENO, buf, 100)) != 0 {
    write(STDOUT, buf, n);
}
```

UNIX KERNEL SUPPORT FOR FILE/I/O



OFT: One entry per file descriptor (file desc flag, ptr to OFT entry)

OFT: File status flag, curr file offset, ptr to V-node table

V-node (i-node table): V-node info, ptr to i-node

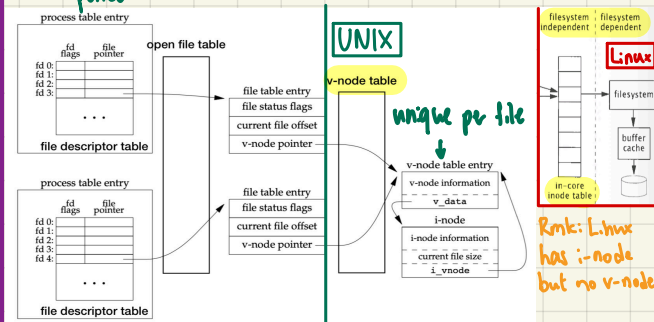
V-node: In-memory data structure for each open file

↳ Info: File type, ptr to func that operate the file

i-node: Stored physically on storage device and in memory

↳ Contains metadata (owner, size, device, protection info, ...)

↳ OS kernel reads i-node from disk to memory when file is opened



FILE I/O CODE SYNTAX (Can end w/ `exit(0);`)

`#include <fcntl.h>` *absolute path => susceptible to TOCTOU attack*
`int open(const char* path, int oflag, ... /* mode_t mode */);`
`int openat(int fd, const char* path, int oflag, ... /* mode_t mode */);` *base of relative path (file dir)*
`int close(int fd);` *If AT_FDCWD, then fd is the curr working dir*
(Rmk: obsolete but `creat(const char path, mode_t mode);` exists, it creates a file if DNE)*

DIFFERENT OFLAGS

Must: `O_RDONLY`, `O_WRONLY`, `O_RDWR` (RD=read, WR=write)

Optional: `O_APPEND`, `O_TRUNC`, `O_CREAT`, `O_NONBLOCK`, `O_SYNC`, `O_DSYNC`, `O_RSYNC`

Mode: `S_I[4 char]` ← order: R(read), W(write), X(execute) ← choose 1 or all three + `USR`, `GRO`, `OTH` (other) ← abbrev to fit length

