


# UNIT 2: UNIX+OS CONCEPTS


## OVERALL STRUCTURE — SYSTEM VS FUNCTION CALLS




	System Call	Function Call
Description	Application calls to functions provided by OS kernel	Program calls to predefined functions (eg: defined in library)
Behavior	Causes switches from user to kernel (space/mode)	Cause no space/mode switches
Speed	Slow	Fast

## PROCESS MANAGEMENT (Time-sharing: Processes can "share" CPU time)

**Context-switch:** The process of storing the execution context of a process; and restoring the execution context of a new process  
**Context-switching causes overload.**



As a flowchart, we have:



## FILE/DIRECTORY MANAGEMENT

In Unix, everything is a file.


Root directory: Top-most dir. "/"

Home directory: Directory when you log in  
 (begin at root)

Absolute path: /x/y/z Relative path: x/y/z

## THE MOUNT COMMAND

Here, we mount in "/mnt",



## SYNTAX

`s=mkdir(name, mode)` create dir  
`s=rmdir(name)` remove dir  
`s=link(name1, name2)` Create new name2=name1  
`s=unlink(name)` remove dir entry  
`s=mount(special, name, flag)` mount  
`s=umount(special)` unmount

Remark:


Pipe ("|") sends the output of one process as the input for another.

# UNIT 3: FILE I/O (PART 1)

## WEEK 2 Shun/335 (@shun4mick)

## FILE I/O (BUFFERED VS UNBUFFERED)

Disk structure:



**Buffered (standard) I/Os:** Functions accumulate results in intermediate buffers, not making system calls each time (e.g. `fread/fwrite`)

**Unbuffered I/Os:** Functions invoke system calls to the kernel each time (e.g. `read()`, `write()` in Unix)

Of course, buffered is faster due to no system calls. Unbuffered would require updating process, OS kernel, and storage as needed. Even with buffer cache in OS kernel, it takes a long time. Besides, buffer cache needs to be maintained carefully alongside storage when `write()` is called.

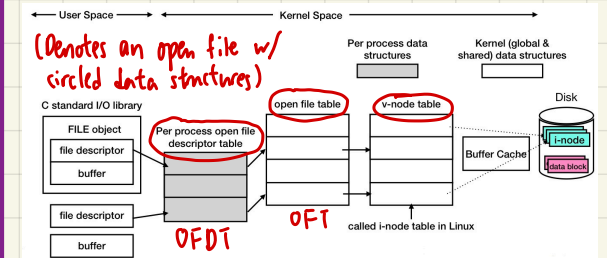
## SYNTAX — FILE DESCRIPTORS

**File descriptor:** A nonnegative integer  $\in [0, \text{OPEN\_MAX}-1]$ , where  $\text{OPEN\_MAX}$  = max files a process can open at once  
 ↳ They are per-process, so diff processes may share file descriptors

In `<unistd.h>`, we have `0=STDIN_FILENO`, `1=STDOUT_FILENO`, `2=STDERR_FILENO` by POSIX.1 standard. More are stored in a file table.

`char buf[100];`  
 E.g. `while (n=read(STDIN_FILENO, buf, 100)) != 0) {`  
`write(STDOUT, buf, n);`

## UNIX KERNEL SUPPORT FOR FILE I/O



**OFT:** One entry per file descriptor (file desc flag, ptr to OFT entry)

**OFT:** File status flag, curr file offset, ptr to V-node table

**V-node i-node table:** V-node info, ptr to i-node

**V-node:** In-memory data structure for each open file

↳ Info: File type, ptr to func that operate the file

**i-node:** Stored physically on storage device and in memory

↳ Contains metadata (owner, size, device, protection info, ...)

↳ OS kernel reads i-node from disk to memory when file is opened

